

CONTENTS

List of Tables	1	
List of Figures	1	
I Introduction	2	
II Problem specification	2	
III Instances	4	
IV Discretizing the configuration space	4	
V Mapping to QUBO	5	
A Binary encoding	5	
B Softening the constraints	5	
VI Solving QUBO instances	5	
A Solving QUBO instances using ICM	5	
VII Quantum Annealing	6	
A Embedding	6	
B Success Probability	6	
VIII Conclusions	6	
IX Acknowledgements	6	
X Figure comments	6	
a Maneuvers	7	
References	7	
		12 (Left) Optimal total delay found by using the Isoenergetic Cluster Method (ICM) at fixed time step Δt , by varying the connected component. Results are for maximum delay time of 60 minutes. (Right) Optimal delay found by using ICM at fixed connected component, by varying the time step Δt 13
		13 . Optimal total delay found by using the Isoenergetic Cluster Method (ICM) at fixed time step Δt as a function of numbers of flight within each connected component. ICM was unable to find solutions for connected component with more than 12 flights. 14
		14 Number of physical qubits versus the number of logical qubits after embedding of QUBO instances for the departure delay model. 14
		15 Average success probability for QUBO instances in dependence of the number of flights N_f and the number of conflicts N_c . The error bars indicate the standard deviation. We used 10000 annealing runs for each instance and penalty weights $\lambda = \lambda_{\text{conflict}} = \lambda_{\text{unique}} \in \{0.5, 1, 2\}$ 15

LIST OF TABLES

I	Parameters of the largest embeddable instances for the D-Wave 2X	6
---	--	---

LIST OF FIGURES

1	Conflict example	3
2	Conflict preprocessing	3
3	Number of connected components vs. d_{max}	9
4	Histogram of connected component sizes	9
5	Histogram of degrees	10
6	Power-law exponent vs. d_{max}	10
7	Histogram of connected component treewidths	11
8	Correlation between connected component size and treewidth	11
9	Treewidth-size correlation coefficient vs. d_{max}	12
10	Effect of discretization on solution quality	12
11	Penalty weight phase diagram	13

Quantum Annealing for Air Traffic Management

Tobias Stollenwerk

German Aerospace Center, Linder Höhe, 51147 Cologne, Germany

Bryan O’Gorman

University of California, Berkeley

*NASA Ames Research Center Quantum Artificial Intelligence Laboratory (QuAIL), Mail Stop 269-1, 94035 Moffett Field CA and
Stinger Ghaffarian Technologies Inc., 7701 Greenbelt Rd., Suite 400, Greenbelt, MD 20770*

Davide Venturelli

*NASA Ames Research Center Quantum Artificial Intelligence Laboratory (QuAIL), Mail Stop 269-1, 94035 Moffett Field CA and
USRA*

Eleanor G. Rieffel

NASA Ames Research Center Quantum Artificial Intelligence Laboratory (QuAIL), Mail Stop 269-1, 94035 Moffett Field CA

Salvatore Mandrà

*NASA Ames Research Center Quantum Artificial Intelligence Laboratory (QuAIL), Mail Stop 269-1, 94035 Moffett Field CA and
Stinger Ghaffarian Technologies Inc., 7701 Greenbelt Rd., Suite 400, Greenbelt, MD 20770*

Olga Rodionova, Hok K. Ng and Banavar Sridhar

tbd

(Dated: May 17, 2017)

In this paper we present the mapping of air traffic management (ATM) problem on quadratic unconstrained boolean optimization (QUBO) problem. After the representation of the ATM problem in terms of a conflict graph, where nodes of the graph represent flights and edges represent a potential conflict between flights, we proceed by discretize the ATM problem and then mapping it in binary variables. As part of our study, we tested the QUBO formulation of the ATM problem using both classical solvers and the D-Wave 2X quantum chip.

I. INTRODUCTION

Efficiently automating air traffic management is increasingly important (increased volume and diversity, environmental concerns, etc.).

Quantum annealing is a promising computational method.

We investigate the feasibility of applying quantum annealing to a particular problem in air traffic management known as “deconflicting”, in which the goal is to modify a set of independently optimal trajectories in a way that removes conflicts between them while minimizing the cost of doing so.

II. PROBLEM SPECIFICATION

The basic input of the deconflicting problem is a set of ideal flight trajectories (space-time paths). These ideal trajectories are specified by the individual flight operators. Each ideal trajectory represents some independent optimization from the operator’s perspective, especially minimizing fuel costs given expected wind conditions between the desired origin and destination at the desired times; for this reason, they are called the “wind-optimal” trajectories. Because of the number of such trajectories and the correlation between them, these trajectories are

likely to conflict; that is, two or more aircraft are likely to get dangerously close to each other if their ideal trajectories are followed without modification. The goal thus is to modify the trajectories to avoid such conflicts.

In theory, the configuration space consists of all physically realistic trajectories; in practice, computational bounds constrain us to consider perturbations of the ideal trajectories. The simplest such perturbation is a departure delay, which is the main focus of the present work. Previous work [1] additionally considered a global perturbation by which a trajectory is sinusoidally shifted parallel to the Earth’s surface. We focus instead on local perturbations to the trajectories, in which a modification to the trajectory is parameterized by some choice of active maneuvers near a potential conflict; such a modification does not affect the preceding part of the trajectory and only affects the subsequent part by the additional delay it introduces.

A full accounting of the cost of such modifications would take into account the cost of departure delays, the change in fuel cost due to perturbing the trajectories, the relative importance of each flight, and many other factors. As in previous work, we consider only the total, unweighted arrival delay, aggregated equally over all of the flights.

Formally, each ideal trajectory $\mathbf{x}_i = (x_{i,t})_{t=t_{i,0}}^{t_{i,1}}$ is specified as a time-discretized path from the departure point $x_{i,t_{i,0}}$ at time $t_{i,0}$ to the arrival point $x_{i,t_{i,1}}$ at time $t_{i,1}$.

For each flight i , the geographical coordinates $x_{i,t}$ (as latitude, longitude, and altitude) are specified at every unit of time (i.e. one minute) between $t_{i,0}$ and $t_{i,1}$; we call this interval $T_i = (t_{i,0}, t_{i,0} + 1, \dots, t_{i,1})$.

For notational simplicity, suppose momentarily that each trajectory \mathbf{x}_i is modified only by introducing delays between time steps. Let $\delta_{i,t}$ be the accumulated delay of flight i at the time that it reaches the point $x_{i,t}$, and let $\delta_{i,t}^*$ be the maximum such delay.

A pair of flights (i, j) are in conflict with each other if any pair of points from their respective trajectories is in conflict. (The trajectories are reasonably assumed to be sufficiently time-resolved so that if the continuously interpolated trajectories conflict then there is a pair of discrete trajectory points that conflict.) A pair of trajectory points $(x_{i,s}, x_{j,t})$ conflict if their spatial and temporal separations are both within the respective mandatory separation standards Δ_x and Δ_t (i.e. 3 nautical miles and 3 minutes):

$$\|x_{i,s} - x_{j,t}\| < \Delta_x \text{ and } |(s + \delta_{i,s}) - (t + \delta_{j,t})| < \Delta_t \quad (1)$$

. The latter condition can be met for some $(\delta_{i,s}, \delta_{j,t}) \in [0, \delta_{i,s}^*] \times [0, \delta_{j,t}^*]$ if and only if

$$\max\{\delta_{i,s}^*, \delta_{j,t}^*\} + \Delta_t > |s - t|, \quad (2)$$

in which case we call the pair of trajectory points *potentially* conflicting. The set C of such pairs of potentially conflicting trajectory points contains strongly correlated clusters. To simplify the constraints, we enumerate all such clusters and refer to them simply as *the* conflicts. That is, we partition the potentially conflicting pairs of trajectory points into disjoint sets,

$$C = \bigcup_k C_k, \quad (3)$$

such that if $\{(i, s), (j, t)\}, \{(i', s'), (j', t')\} \in C_k$ for some k then $i = i' < j = j'$ and for all $s'' \in [\min\{s, s'\}, \max\{s, s'\}]$ there exists some $t'' \in [\min\{t, t'\}, \max\{t, t'\}]$ such that $\{(i, s''), (j, t'')\} \in C_k$. Thus every conflict k is associated with a pair of flights $I_k = \{i, j\}$. Let $K_i = \{k | i \in I_k\}$ be the set of conflicts to which flight i is associated.

Having identified disjoint sets of conflicts, we relax the supposition that the trajectory modifications only introduce delays between time steps. Instead, we consider modifications to the trajectories that introduce delays local to particular conflicts. Specifically, the configuration space consists of the departure delays $\mathbf{d} = (d_i)_{i=1}^n$ and the set of local maneuvers $\mathbf{a}_k = (\mathbf{a}_k)_k$, where \mathbf{a}_k represents some parameterization of the local maneuvers used to avoid conflict k . Let $d_{i,k}(\mathbf{d}, \mathbf{a}_k)$ be the delay introduced to flight i at conflict k , as a function of the departure delays and local maneuvers. With this notation, we can write the total delay as

$$D = \sum_{i=1}^n \left(d_i + \sum_{k \in K_i} d_{i,k} \right). \quad (4)$$

This is the quantity we wish to minimize subject to avoiding all potential conflicts.




FIG. 1. Example of two parallel potential conflicts between two transatlantic flights starting from the east coast of the USA.




FIG. 2. Preprocessing: Reduction in the number of potential conflicts for various upper delay bounds D_{\max} .

A conflict can be avoided locally by introducing earlier delays differentially, thereby increasing the temporal separation; by some active maneuver of one or both of the flights; or by some combination thereof. We focus on the former case. Let

$$D_{i,k} = d_i + \sum_{k' \in K_i | k' < k} d_{i,k'} \quad (5)$$

be the accumulated delay of flight i by the time it reaches conflict k . We assume that the set of conflicts K_i associated with flight i is indexed in temporal order, i.e. if $k' < k$ and $k, k' \in K_i$, then flight i reaches conflict k' before conflict k . The pairs of conflicting trajectory points associated with conflict k are given by

$$T_k = \{(s, t) | \{(i, s), (j, t)\} \in C_k, i < j\}. \quad (6)$$

Thus the potential conflict is avoided only if

$$D_{i,k} - D_{j,k} \notin D_k \quad (7)$$

where

$$D_k = \bigcup_{(s,t) \in T_k} (-\Delta_t + t - s, \Delta_t + t - s) = [\Delta_k^{\min}, \Delta_k^{\max}], \quad (8)$$

$$\Delta_k^{\min} = 1 - \Delta_t + \min_{(s,t) \in T_k} \{t - s\}, \quad (9)$$

$$\Delta_k^{\max} = \Delta_t - 1 + \max_{(s,t) \in T_k} \{t - s\}. \quad (10)$$

In the remainder of this paper, we focus on the restricted problem in which only departure delays are allowed. In this simplified case, the configuration space is simply $\mathbf{d} = (d_i)_{i=1}^n$, the cost function simply $D = \sum_{i=1}^n d_i$, and the constraints simply $d_i - d_j \notin D_k$ for all k .

III. INSTANCES

To assess our methods on realistic instances of the problem, we use the actual wind-optimal trajectories for transatlantic flights on July 29, 2012, as was done in previous work [1]. In these trajectories, each flight has a constant (cruising) altitude and constant speed, to within (classical) machine precision, though our methods generalize to instances without these special properties.

One perspective into the nature of an instance of the deconflicting problem is the *conflict graph*, whose vertices correspond to flights and which has an edge between a pair of vertices if there is at least one potential conflict between the corresponding flights. Note that the conflict graph for a given set of trajectories depends on the parameters of the problem. In the case of only departure delays, whether or not a potential conflict, and thus an edge in the conflict graph, exists between two flights is a function of d_{\max} . For a certain value of d_{\max} , the conflict graph may contain several connected components, which can be considered as smaller, independent instances. Figure 3 shows this dependence of the number of connected components on the maximum delay d_{\max} , and Figure 4 shows the distribution of the sizes of the connected components for various values of d_{\max} . Interestingly, most of the connected components are very small; for example, with $d_{\max} = 60$ minutes, approximately 75% of the connected components contain no more than 10 flights.

As part of our analysis, we also studied the probability distribution of the connectivity, namely the number of flights for which a given flight share a potential conflict with. Figure (5) shows the distribution of degrees in the conflict graph for $d_{\max} = 60$, which seem to be approximately distributed according to a power law, i.e. the number of vertices with degree d is proportional to $d^{-\alpha}$. This is consistent with a so-called “small-world” model believed to be typical of many real-world graphs [2], which are generated by preferential attachment and resultingly contain a few number of highly-connected hubs, as is the case with air traffic. Figure (6) shows the dependence of this empirical power-law exponent α as a function of d_{\max} . As d_{\max} increases, the exponent decreases. The

larger the delay, the less the structure of the trajectories matters and the flatter the distribution of degrees in the conflict graph.

In many cases, generally hard problems are easy when restricted to tree-like instances [3, 4]. For example, if the conflict graph here is a tree, then the optimum could be easily found by propagating the delays along the tree; on the other hand, if the conflict graph is a complete graph, finding the optimum is much harder. The tree-width of a graph formalizes this notion of tree-likeness, ranging from 1 for a tree to $n - 1$ for fully connected graph. We examine the treewidth of the connected components as a proxy for the hardness of the instances they represent. Figure 7 shows that most connected components are very tree-like, while there are a few strongly-connected components.

Figure 8 shows that the treewidth of a connected component scales approximately linearly with its size. This indicates that realistic instances of the deconflicting are indeed hard, and not restricted to easier (bounded tree-width) instances of the generally hard problem. Moreover, the correlation γ between the tree-width of a connected component and its size increases with d_{\max} , as shown in Figure 9. The larger d_{\max} , the more potential conflicts there are; restricting d_{\max} also restricts the number of conflicts.

IV. DISCRETIZING THE CONFIGURATION SPACE

To apply quantum annealing to the deconflicting problem, we must encode the configuration space \mathbf{d} in binary-valued variables. To do so, we must first discretize and bound the allowed values. Let Δ_d be the resolution of the allowed delays and $d_{\max} = N_d \Delta_d$ the maximum allowed delay, so that $d_i \in \{\Delta_d l | l \in [0, 1, \dots, N_d]\}$. The larger the configuration space, the more qubits needed to encode it, and so determining the effect of this discretization on solution quality is integral to the effective use of quantum annealing. To do so, we solve the deconflicting problem with departure delays only for various delay resolutions and upper bounds and compare the various optima to the continuous problem without restrictions (other than non-negativity) on the delays. The exact optima are found by modeling the problem as a constraint satisfaction problem [5]. We used as problem instances the non-trivially sized connected components of the conflict graph (as described in the previous section) for $d_{\max} = 18$; the largest such instance has 50 flights and 104 potential conflicts.

Figure 10 shows the minimum total delay of a problem instance with 19 flights and 47 potential conflicts for various values of Δ_d and d_{\max} . With the exception of the small maximum delay $d_{\max} = 3$, the total delay of the solutions is nearly independent of the maximum delay. The total delay is non-decreasing with respect to the coarseness Δ_d of the discretization for a fixed maximum delay d_{\max} , and non-increasing with respect to d_{\max} for a fixed Δ_d . Since the original data is discretized in time in units

of 1 minute, $\Delta_d = 1$ yield the same result as a continuous variable with the same upper bound. Above some threshold value d_{\max}^0 , further increasing the maximum delay does not decrease the minimum total delay. With one exception, we found that for all the investigated problem instances $d_{\max}^0 \leq 6$ minutes (see Figure 10). Therefore we conclude, that a moderate maximum delay is sufficient even for larger problem instances. On the other hand, the delay discretization should be as fine as possible to obtain a high quality solutions.

V. MAPPING TO QUBO

A. Binary encoding

Having suitably discretized the configuration space, we must then encode it into binary-valued variables. The value of d_i is encoded in $N_d + 1$ variables $d_{i,0}, \dots, d_{i,N_d+1} \in \{0, 1\}$ using a one-hot encoding:

$$d_{i,\alpha} = \begin{cases} 1, & d_i = \alpha, \\ 0, & d_i \neq \alpha; \end{cases} \quad d_i = \Delta_d \sum_{l=0}^{N_d} d_{i,l}. \quad (11)$$

To enforce this encoding, we add the penalty function

$$f_{\text{encoding}} = \lambda_{\text{encoding}} \sum_{i=1}^n \left(\sum_{l=0}^{N_d} d_{i,l} - 1 \right)^2, \quad (12)$$

where $\lambda_{\text{encoding}}$ is a penalty weight sufficiently large to ensure that any cost minimizing state satisfies $f_{\text{encoding}} = 0$. In terms of these binary variables, the cost function is

$$f_{\text{delay}} = \Delta_d \sum_{i=1}^n \sum_{l=0}^{N_d} d_{i,l}, \quad (13)$$

Lastly, actualized conflicts are penalized by

$$f_{\text{conflict}} = \lambda_{\text{conflict}} \sum_k \sum_{\substack{l, l' | \Delta_d(l-l') \in D_k \\ i, j \in I_k | i < j}} d_{i,l} d_{j,l'}, \quad (14)$$

where again $\lambda_{\text{conflict}}$ is a sufficiently large penalty weight. The overall cost function to be minimized is

$$f = f_{\text{encoding}} + f_{\text{delay}} + f_{\text{conflict}}. \quad (15)$$

B. Softening the constraints

In the QUBO formalism, there are no hard constraints; thus the use of penalty functions in the previous section. For sufficiently large penalty weights, the optimum will satisfy the desired constraints. However, precision is a limited resource in quantum annealing [?]; therefore, we would like to determine the smallest sufficient penalty weights.

For a given instance, we say that a pair of penalty weights $(\lambda_{\text{conflict}}, \lambda_{\text{encoding}})$ is valid if the minimum of the total cost function satisfies both the conflict and encoding constraints when using those weights. Figure 11 shows the phase space of these penalty weights for a single instance with 7 flights and 9 conflicts. The box-like boundary between valid and invalid penalty weights suggests that the validity of the two penalty weights is independent; this box-like boundary is found for all of our instances with up to 7 flights and 9 conflicts.

One can give an upper bound for the sufficiently large penalty weights by the following considerations. A minimal violation of the hard constraints yield an additional contribution to the QUBO cost function of λ_{unique} or $\lambda_{\text{conflict}}$, respectively. Such a violation would correspond to single bit flip in the binary delay variables $d_{i,\alpha}$. Therefore the contribution from (??) would be reduced maximally by

$$\min_{\alpha} \frac{-\alpha}{d_{\max}} = -1$$

Hence, sufficiently large penalty weights must fulfill the following conditions

$$\begin{aligned} \lambda_{\text{unique}} &> 1 \\ \lambda_{\text{conflict}} &> 1. \end{aligned}$$

This corresponds to a box like shape as it appears in figure 11.

VI. SOLVING QUBO INSTANCES

A. Solving QUBO instances using ICM

In this section we present the results for the optimization of the QUBO formulation of the ATM problem using the Isoenergetic Cluster Method (a rejection-free cluster algorithm for spin glasses that greatly improves thermalization) [6], which has been shown to be one of the fastest classical heuristic to optimize QUBO problems [7].

Figure (12) shows the total delay time optimized by ICM either by varying the partition at fixed the delay step Δt (left panel) or by varying the delay step Δt at fixed partition (right panel). As one can see, the total delay decreases by decreasing Δt and it eventually reaches an optimal plateau. Results are for maximum delay of 60 minutes. This is consistent with the idea that smaller Δt allows a finer optimization of the delays of the flights.

In Figure (13) we show the optimal delay time found by ICM as a function of the number of the flights in the connected components. Results are for a maximum delay of 60 minutes. Unfortunately, ICM was unable to optimize connected components with more than 12 flights.

This can be explained by recalling that ICM works the best for almost-planar problem while its performance quickly decreases for fully-connected problems. Indeed, as shown in Section III, the underlying graph of connected components look more like a fully-connected graph rather than a tree graph by increasing the number of flights inside the connected component.

VII. QUANTUM ANNEALING

In this section we report on our efforts to solve problem instances from the departure delay model from Section ?? with a D-Wave 2X quantum annealer. We restricted ourselves to instances with $d_{\max} = D_{\max} = 18$ and $\Delta_d \in \{3, 6, 9\}$.

A. Embedding

In order to make a QUBO amenable for a D-Wave 2X quantum annealer, it has to obey certain hardware constraints. For instance the connections between the binary variables are restricted to the so called Chimera graph [8]. However, it is possible to map every QUBO to another QUBO which obeys the Chimera architecture while increasing the number of binary variables used by a so called minor-embedding technique [9].

Δ_d	3	6	9
Number of flights N_f	13	19	50
Number of conflicts N_c	27	47	104
Number of logical qubits	91	76	150
Average number of physical qubits	631	395	543

TABLE I. Parameters of the largest embeddable instances for the D-Wave 2X

We found, that non of the instances were suitable for direct calculation on the D-Wave machine. Therefore we used D-Wave’s heuristic embedding algorithm[10] to embed instances with up to $N_f = 50$ and $N_c = 104$ depending on discretization (cf. Table I). We used up to 5 different embeddings for each QUBO instance. In figure 14 on can see the dependence of the number of physical qubits on the number of logical qubits.

B. Success Probability

In order to investigate the performance of the D-Wave 2X machine, we compared its results to the ones of an exact solver. We used an exact Max-SAT solver [11] after we mapped the QUBOs to Max-SAT [12]. For each QUBO instance, we ran the annealing process $N_r = 10000$ times. The success probability is then given by the ratio of the number of annealing solutions which are equal to exact solution and the N_r .

In figure 15 the dependence of the average success probability on the number of flights and the number of conflicts is shown. The average was taken over all successfully embedded QUBO instances with the same number of flights and number of conflicts, respectively. On can see, that the success probability decreases for larger problem instances as well as for finer discretizations. This is a result of the limited precision in the specification of a QUBO on the D-Wave 2X machine [13].

VIII. CONCLUSIONS

TODO

IX. ACKNOWLEDGEMENTS

X. FIGURE COMMENTS

1. Figure 2
 - (a) This can probably be removed. Essentially the same information as Figure 4.
2. Figure 3
 - (a) Why is this not monotonic?
3. Figure 4
 - (a) At the very least, we should get rid of all points that are 0 to make clear where there are points that aren’t.
 - (b) What does “regardless of the maximum delay time” mean? The union over various delay times? If so, we should at least specify which delay times, but that doesn’t seem to be a very meaningful number. It would be much better to have different curves for particular values of d_{\max} .
4. Figure 6.
 - (a) How exactly is α calculated?
5. Figure 7
 - (a) Would be much better if for various values of d_{\max} .
 - (b) Does this figure add much given Figure 4?
6. Figure 10
 - (a) Top: can remove minutes
 - (b) Bottom: remove d_{\max}^0 .
 - (c) Bottom: how were the instances chosen?

a. Maneuvers

A more realistic model of the problem can be created by including maneuvers. As mentioned above the maneuvers enter our formulation as additional delays d_{ik} at the conflict time. In the course of mapping to a QUBO formulation, we need to make sure to retain the combinatorial nature of the problem. We do this by restricting the vast realm of maneuvers to two distinct choices: Only one of the two involved flights is delayed while leaving the other flight untouched

$$\text{if } d_{ik} \neq 0 \Rightarrow d_{jk} = 0 \quad \forall (i, j) \in I_k \quad \forall k. \quad (16)$$

Moreover, we set the resulting maneuver delays to a constant value d_M large enough to capture all kinds of real maneuvers. A natural choice for this is the temporal conflict threshold $d_M = \Delta_t$.

With (??) we can introduce the delay a flight i at the conflict k as

$$D_{ik} = d_i + \sum_{k' < k} d_{ik'}, \quad (17)$$

where we have defined a temporal ordering of the conflicts for each flight i by

$$\begin{aligned} k < p & \text{ if } t < t' \\ \text{for } t &= \min_s x_{i,s} \in C_k, \\ t' &= \min_s x_{i,s} \in C_p \end{aligned}$$

The departure delay variables are represented by binary variables as it was done in Section ???. The maneuver delays are given by

$$d_{ik} = d_M a_{ik} \quad a_{ik} \in \{0, 1\}$$

Since the total delay is given by $\sum_{ik} D_{ik}$, we can write the corresponding QUBO contribution as

$$\tilde{Q}_{\text{delay}} = \sum_{i\alpha} \alpha d_{i\alpha} + \sum_{ik} d_M a_{ik},$$

For the conflict avoidance, we need to introduce another variable representing the delay at a given conflict

$$D_{ik} = \sum_{\delta} \delta \Delta_{ik\delta} \quad \Delta_{ik\delta} \in \{0, 1\}.$$

By restricting ourselves to $\Delta_d = \Delta_t$ the values of δ in the above equation are given as

$$\delta \in \{0, \Delta_t, 2\Delta_t, \dots, (N_d + M_{ik})\Delta_t\}.$$

Here, M_{ik} is the number of conflicts the flight i is involved in before k . In order to fulfill (17) we add the following contribution to the QUBO

$$\tilde{Q}_{\Delta} = \lambda_{\Delta} \sum_{ik} \left(\sum_{\alpha} \alpha d_{i\alpha} + \sum_{k' < k} d_M a_{ik'} - \sum_{\delta} \delta \Delta_{ik\delta} \right)^2 \Big|_{i,j \in I_k}$$

For unique representation of the variables we add

$$\begin{aligned} \tilde{Q}_{\text{unique}} = \lambda_{\text{unique}} \left\{ \sum_i \left(\sum_{\alpha} d_{i\alpha} - 1 \right)^2 \right. \\ \left. + \sum_{ik} \left(\sum_{\delta} \Delta_{ik\delta} - 1 \right)^2 \right\}. \end{aligned}$$

Conflicts are avoided if $D_{ik} - D_{jk} \notin D_k$, $(i, j) \in I_k$. The corresponding QUBO contribution reads

$$\tilde{Q}_{\text{conflict}} = \lambda_{\text{conflict}} \sum_k \sum_{(\delta, \delta') \in B_k} \Delta_{ik\delta} \Delta_{jk\delta'} \Big|_{i,j \in I_k}$$

where B_k is the set of all (δ, δ') which correspond to a conflict

$$B_k = \{(\delta, \delta') \mid \delta - \delta' \in D_k\}$$

The penalty weights λ_{Δ} , λ_{unique} and $\lambda_{\text{conflict}}$ must be chosen large enough to ensure vanishing contributions from the corresponding QUBO terms for the solution.

Finally, the maneuver decision described by (16) is incorporated by a antiferromagnetic coupling between the two maneuver delay variables

$$\tilde{Q}_{\text{maneuver}} = J \sum_k (s_{ik} s_{jk} + 1)_{i,j \in I_k}.$$

with

$$s_{ik} = 2a_{ik} - 1 \in \{-1, 1\}$$

and $J > 0$ has to be chosen large enough. A solution is considered to be valid only if $\tilde{Q}_{\text{maneuver}} = 0$. Hence, the total QUBO for the maneuver model reads

$$Q_{\text{MM}} = \tilde{Q}_{\text{delay}} + \tilde{Q}_{\Delta} + \tilde{Q}_{\text{unique}} + \tilde{Q}_{\text{conflict}} + \tilde{Q}_{\text{maneuver}}$$

- Conference (2016).
- [2] R. Albert, H. Jeong, and A.-L. Barabási, *Internet: Diameter of the world-wide web*, nature **401**, 130 (1999).
 - [3] U. Bertele and F. Brioschi, *Nonserial dynamic programming* (Academic Press, 1972).
 - [4] R. Halin, *S-functions for graphs*, Journal of geometry **8**, 171 (1976).
 - [5] E. Hebrard, E. O'Mahony, and B. O'Sullivan, *Constraint Programming and Combinatorial Optimisation in Numberjack* (Springer Berlin Heidelberg, Berlin, Heidelberg, 2010), pp. 181–185, ISBN 978-3-642-13520-0, URL http://dx.doi.org/10.1007/978-3-642-13520-0_22.
 - [6] Z. Zhu, A. J. Ochoa, and H. G. Katzgraber, *Efficient cluster algorithm for spin glasses in any space dimension*, Physical review letters **115**, 077201 (2015).
 - [7] S. Mandrà, Z. Zhu, W. Wang, A. Perdomo-Ortiz, and H. G. Katzgraber, *Strengths and weaknesses of weak-strong cluster problems: A detailed overview of state-of-the-art classical heuristics versus quantum approaches*, Physical Review A **94**, 022337 (2016).
 - [8] Note1, reference to Chimera graph.
 - [9] Note2, reference to minor-embedding.
 - [10] Note3, reference to D-Wave embedding algorithm.
 - [11] Note4, reference to akmaxsat.
 - [12] Note5, reference on mapping QUBO to Max-SAT.
 - [13] Note6, reference to limited precision of D-Wave.

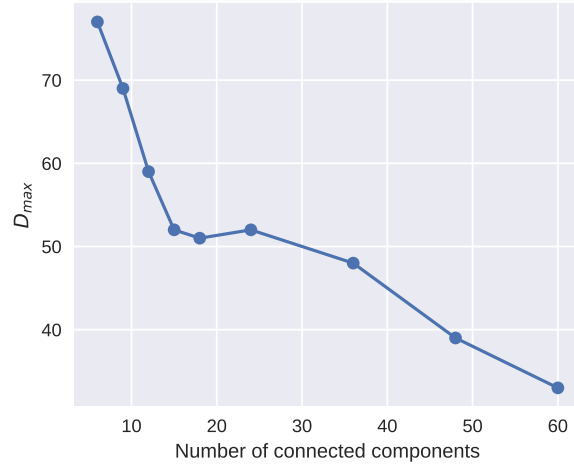


FIG. 3. Number of connected components versus d_{max} .




FIG. 4. Histogram of the number of flights per connected component, regardless the maximum delay time.




FIG. 5. Histogram of the degrees of vertices in the conflict graph for $d_{\max} = 60$. The distribution of the degrees approximately follows a power law, with the exponent depending on d_{\max} .




FIG. 6. Empirical power-law exponent versus d_{\max} .

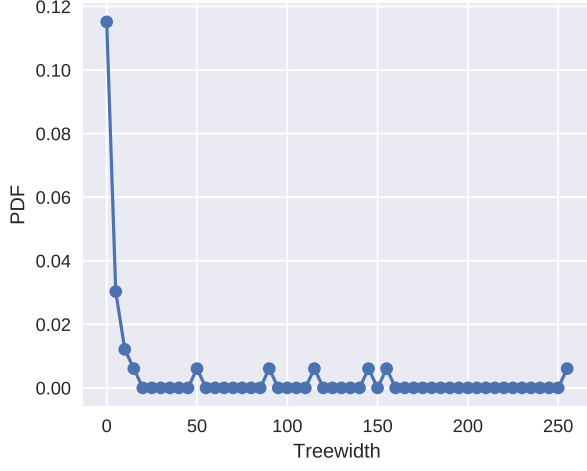


FIG. 7. Histogram of the treewidths of the connected components of the conflict graph for various values of d_{\max} .

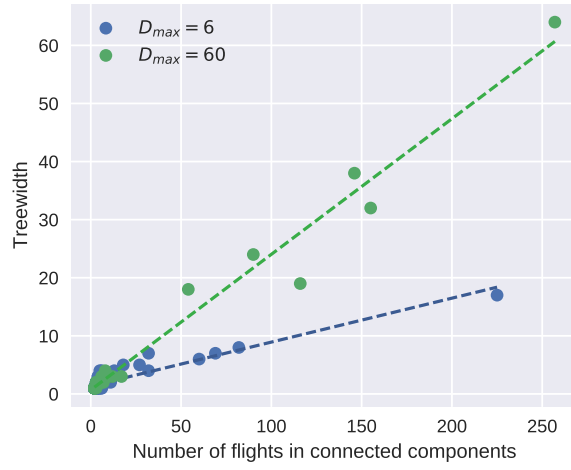


FIG. 8. The treewidths of connected components versus their sizes for various values of d_{\max} . The correlation is approximately linear, with a slope γ that depends on d_{\max} .




FIG. 9. Slope γ as a function of the maximum delay time.

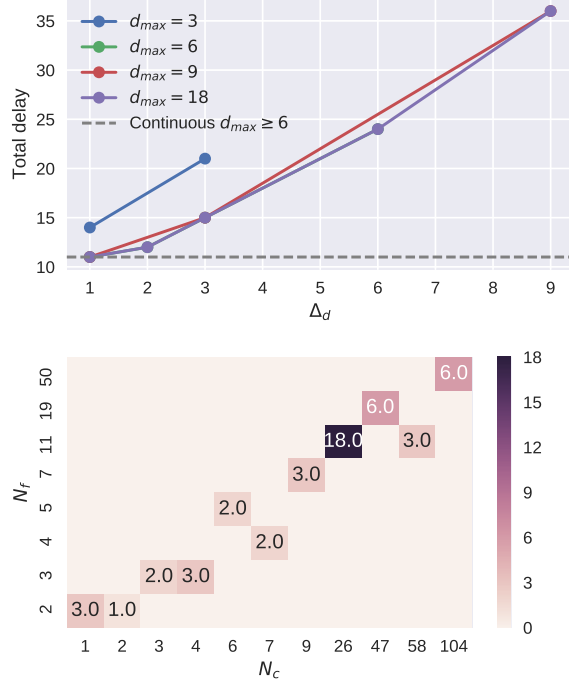


FIG. 10. Top: Minimum total delay of a problem instance with 19 flights and 47 conflicts for various values of Δ_d and d_{\max} . Bottom: Minimum d_{\max} necessary to obtain same optimum as that without bounding the delay, for all non-trivial instances with $D_{\max} = 18$ and up to 50 flights and 104 conflicts.




FIG. 11. Validity of exact solution to a QUBO extracted from a problem instance with $N_f = 7$ flights and $N_c = 9$ conflicts in dependence on the choice of the penalty weights, λ_{unique} and $\lambda_{\text{conflict}}$. Here, $\Delta t = 3$ and $d_{\text{max}} = 18$.





FIG. 12. (Left) Optimal total delay found by using the Isoenergetic Cluster Method (ICM) at fixed time step Δt , by varying the connected component. Results are for maximum delay time of 60 minutes. (Right) Optimal delay found by using ICM at fixed connected component, by varying the time step Δt .




FIG. 13. . Optimal total delay found by using the Isoenergetic Cluster Method (ICM) at fixed time step Δt as a function of numbers of flight within each connected component. ICM was unable to find solutions for connected component with more than 12 flights.




FIG. 14. Number of physical qubits versus the number of logical qubits after embedding of QUBO instances for the departure delay model.

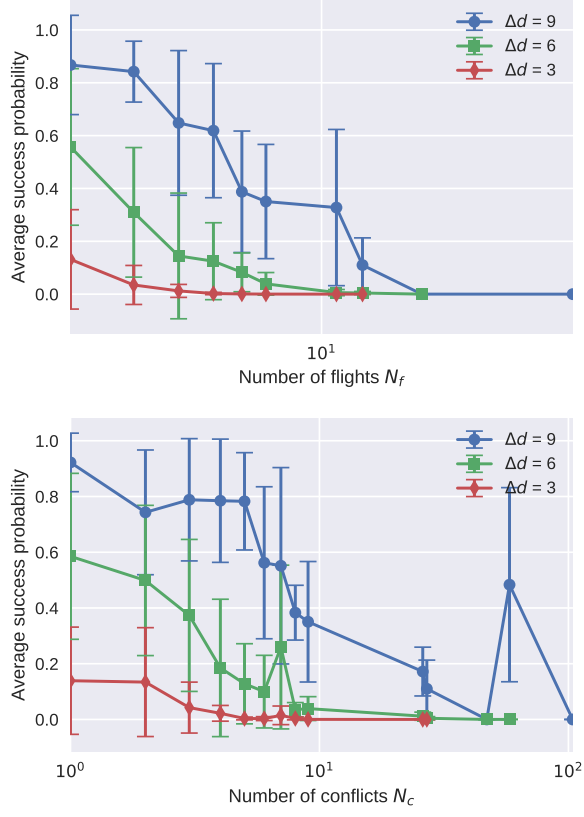


FIG. 15. Average success probability for QUBO instances in dependence of the number of flights N_f and the number of conflicts N_c . The error bars indicate the standard deviation. We used 10000 annealing runs for each instance and penalty weights $\lambda = \lambda_{\text{conflict}} = \lambda_{\text{unique}} \in \{0.5, 1, 2\}$.